

STAT 123 - Homework 3

Parker DeBruyne - V00837207

09/03/2022

1. Download and save the homework3Data.csv dataset and read it into R. This data set contains 6 numerical variables A, B, C, X, Y, Z.

```
hw3Data = read.csv("homework3Data.csv")
head(hw3Data, 3)
```

```
##   X.1      A      B      C X      Y      Z
## 1   1 0.9102545 1.106865 124.6358 9 -0.7828927 0.9983816
## 2   2 1.0540369 2.104931 214.5555 19 -0.1010444 1.2075391
## 3   3 0.6864621 3.361089 170.2348 14  1.9741031 1.4268379
```

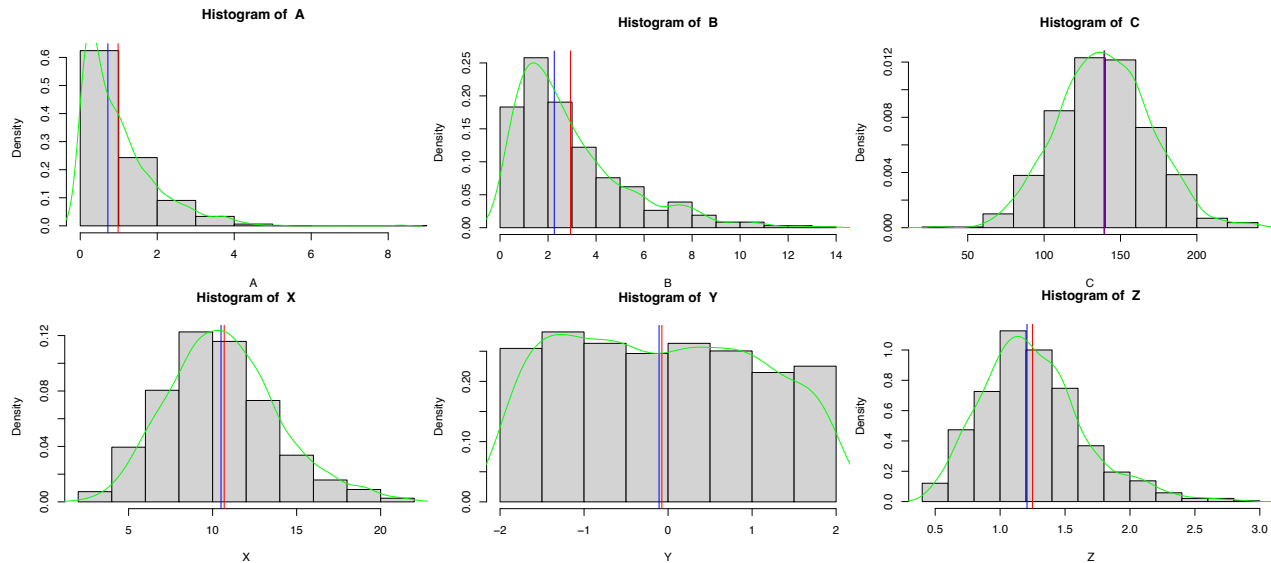
(a) If you use the function `hist()` to plot a histogram followed by the function `abline(v=3,col='red')`, this will add a red vertical line at $x = 3$.

Using these functions:

- plot a histogram for each of the variables.
- add vertical lines for the sample mean and the sample median of those variables. Make the sample mean lines red and the sample median lines blue.
- add a green density curve to each plot.
- make sure your histogram has a main title.

```
hw3Data_names = colnames(hw3Data)

for(i in 2:length(hw3Data_names)){
  hist(hw3Data[[i]],
      main = paste("Histogram of ",
                  hw3Data_names[i]),
      prob=TRUE,
      xlab = hw3Data_names[i])
  abline(v=mean(hw3Data[[i]]),col='red')
  abline(v=median(hw3Data[[i]]),col='blue')
  lines(density(hw3Data[[i]]),col='green')
}
```



(b) One of the variables is normally distributed. Determine which variable it is and justify why you think it is that variable.

Ans: Variable C is normally distributed because the mean and the median are approximately equal and overlap.

(c) For the normally distributed variable you identified in part (b), use the 68 – 95 – 99.7 rule to determine the intervals such that approximately 68% of the data, 95% of the data, and 99.7% of the data lie within those intervals.

```
sigma_C = sd(hw3Data[["C"]])
mean_C = mean(hw3Data[["C"]])

C_68 = round((mean_C + sigma_C), 2)
print(paste("68th percentile =", C_68))

## [1] "68th percentile = 169.94"

C_95 = round((mean_C + 2*sigma_C), 2)
print(paste("95th percentile =", C_95))

## [1] "95th percentile = 200.27"

C_99.7 = round((mean_C + 3*sigma_C), 2)
print(paste("99.7th percentile =", C_99.7))

## [1] "99.7th percentile = 230.59"
```

(d) Use the quantile() function to approximate those same intervals. Are the intervals the same?

```
quantile(hw3Data[["C"]], 0.68)

##      68%
## 153.7819
```

```
quantile(hw3Data[["C"]], 0.95)
```

```
##      95%  
## 190.4729
```

```
quantile(hw3Data[["C"]], 0.997)
```

```
##      99.7%  
## 232.3657
```

(e) Use the `qnorm()` function (with the sample mean and sample standard deviation) to approximate those same intervals. Are these intervals the same as the intervals in either part (c) or part (d)?

```
round(qnorm(c(0.68,0.95,0.997), mean= mean_C, sd=sigma_C), 2)
```

```
## [1] 153.80 189.50 222.94
```

Ans: They're very close, but not exactly the same!

(f) Suppose you wish to estimate the population mean for the normally distributed variable you identified in part (b). Compute the following:

- an estimate of the population mean.
- the estimated standard error of the statistic.
- the critical value for an 88% confidence interval.
- a 88% confidence interval for the population mean.

```
n = length(hw3Data)
```

```
#Estimate the population mean
```

```
boot_sample = numeric()
```

```
for(i in 1:10000){
```

```
  temp_samp = sample(hw3Data[["C"]], n, replace=TRUE)
```

```
  temp_mean = mean(temp_samp)
```

```
  boot_sample[i] = temp_mean
```

```
}
```

```
boot_mean = mean(boot_sample)
```

```
print(paste("Bootstrapped population mean =", round(boot_mean, 2)))
```

```
## [1] "Bootstrapped population mean = 139.51"
```

```
#Estimate the standard error
```

```
std_err = sd(hw3Data[["C"]])/sqrt(n)
```

```
print(paste("Standard error of C =", round(std_err,2)))
```

```
## [1] "Standard error of C = 11.46"
```

```
#Estimate the critical value for an 88% confidence interval.
```

```
diff_88 = (100-88)/2
```

```
interval_88 = c(diff_88, 100-diff_88)
```

```
crit_val_88 = qnorm(interval_88/100, mean=0, sd=1)
```

```
print(paste("Critical Value estimate for an 88% confidence interval =", round(crit_val_88[2], 2)))
```

```
## [1] "Critical Value estimate for an 88% confidence interval = 1.55"
#Estimate a 88% confidence interval for the population mean.

confidence_int_88 = c(boot_mean - crit_val_88[2]*std_err, boot_mean + crit_val_88[2]*std_err)

print(paste("Confidence interval for 88% = (",
            round(confidence_int_88[1], 2),
            ",",
            round(confidence_int_88[2], 2),
            ")"))

## [1] "Confidence interval for 88% = ( 121.69 , 157.33 )"
```

2. For this question, you will need to install the package ‘dplyr’ into R by typing in the command `install.packages('dplyr')`. Then you need to load dplyr into R by typing in the command `library(dplyr)`. We will be using the starwars data set that is built into the dplyr package.

```
library(dplyr)

##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##     filter, lag
## The following objects are masked from 'package:base':
##
##     intersect, setdiff, setequal, union
```

(a) Create a vector called `names` which contains the names of starwars characters that are included in the data set.

```
names = starwars$name
```

(b) The function `nchar()` determines the number of characters in a string. How many characters are in the 5th, 20th, and 34th elements of the `names` vector?

```
print(paste("Number of characters in the 5th element = ",nchar(names[5])))

## [1] "Number of characters in the 5th element = 11"
print(paste("Number of characters in the 20th element = ",nchar(names[20])))

## [1] "Number of characters in the 20th element = 9"
print(paste("Number of characters in the 34th element = ",nchar(names[34])))

## [1] "Number of characters in the 34th element = 13"
```

(c) Create an empty numeric vector called `num_char`. Write a loop which calculates the number of characters in each element of the `names` vector, and puts the corresponding number in the `num char` vector.

```
num_char = numeric()
for(i in 1:length(names)){
  num_char[i] = nchar(names[i])
}
```

```
num_char
```

```
## [1] 14  5  5 11 11  9 18  5 17 14 16 14  9  8  6 21 14 16  4  9  9  5  5 16  5
## [26]  6 10 12 21  9 12 11 13 13 12 10  8  5  7 13 14 10 11 11  8  7 14 10 12  9
## [51]  9 10 11 11  8 10 12  5 11 17 15 13  5  5 19 10 10 15  7  7 10 13  6 10  8
## [76]  8  8  7 15  9 10  4  3 11  3 14 13
```

(d) Now do the same thing that you did in part (c) using the `lapply()` or `sapply()` function in R. Be careful that your output is a vector.

```
num_char_sapply = sapply(names, nchar)
num_char_sapply
```

```
##      Luke Skywalker      C-3P0      R2-D2
##           14           5           5
##      Darth Vader      Leia Organa      Owen Lars
##           11           11           9
##      Beru Whitesun lars      R5-D4      Biggs Darklighter
##           18           5           17
##      Obi-Wan Kenobi      Anakin Skywalker      Wilhuff Tarkin
##           14           16           14
##           Chewbacca      Han Solo      Greedo
##           9           8           6
##      Jabba Desilijic Tiure      Wedge Antilles      Jek Tono Porkins
##           21           14           16
##           Yoda      Palpatine      Boba Fett
##           4           9           9
##           IG-88      Bossk      Lando Calrissian
##           5           5           16
##           Lobot      Ackbar      Mon Mothma
##           5           6           10
##      Arvel Crynyd Wicket Systri Warrick      Nien Nunb
##           12           21           9
##      Qui-Gon Jinn      Nute Gunray      Finis Valorum
##           12           11           13
##      Jar Jar Binks      Roos Tarpals      Rugor Nass
##           13           12           10
##           Ric Olié      Watto      Sebulba
##           8           5           7
##      Quarsh Panaka      Shmi Skywalker      Darth Maul
##           13           14           10
##      Bib Fortuna      Ayla Secura      Dud Bolt
##           11           11           8
##           Gasgano      Ben Quadinaros      Mace Windu
##           7           14           10
```

##	Ki-Adi-Mundi	Kit Fisto	Eeth Koth
##	12	9	9
##	Adi Gallia	Saesee Tiin	Yarael Poof
##	10	11	11
##	Plo Koon	Mas Amedda	Gregar Typho
##	8	10	12
##	Cordé	Cliegg Lars	Poggle the Lesser
##	5	11	17
##	Luminara Unduli	Barriss Offee	Dormé
##	15	13	5
##	Dooku	Bail Prestor Organa	Jango Fett
##	5	19	10
##	Zam Wesell	Dexter Jettster	Lama Su
##	10	15	7
##	Taun We	Jocasta Nu	Ratts Tyerell
##	7	10	13
##	R4-P17	Wat Tambor	San Hill
##	6	10	8
##	Shaak Ti	Grievous	Tarfful
##	8	8	7
##	Raymus Antilles	Sly Moore	Tion Medon
##	15	9	10
##	Finn	Rey	Poe Dameron
##	4	3	11
##	BB8	Captain Phasma	Padmé Amidala
##	3	14	13

3. Consider again the homework3Data.csv dataset and the variable X.

(a) Write a bootstrap computing the median on 10, 000 samples (with replacement) of size 600 of the variable X.

```
n = length(hw3Data[["X"]])

boot_sample = numeric()
for(i in 1:10000){
  temp_samp = sample(hw3Data[["X"]], 600, replace=TRUE)
  temp_median = median(temp_samp)
  boot_sample[i] = temp_median
}


boot_median = mean(boot_sample)
```

(b) Plot the resulting sampling distribution for the median of X.

```
hist(boot_sample,
     main="Distribution of Bootstrapped Median for X",
     xlab="Median of X",
     prob=TRUE)
abline(v=mean(boot_sample),col='red')
abline(v=median(boot_sample),col='blue')
```

```
lines(density(boot_sample),col='green')
```

Distribution of Bootstrapped Median for X



(c) Determine an estimate for the median of X.

```
boot_median = mean(boot_sample)
print(paste("Estimated median for the population of X =", round(boot_median, 2)))
```

```
## [1] "Estimated median for the population of X = 10.49"
```

(d) Compute a 95% confidence interval for the median of X.

```
#Estimate the standard error
std_err = sd(boot_sample)/sqrt(n)

#Estimate the critical value for an 95% confidence interval.
diff_95 = (100-95)/2
interval_95 = c(diff_95, 100-diff_95)
crit_val_95 = qnorm(interval_95/100, mean=0, sd=1)

#Estimate a 95% confidence interval for the population median.
confidence_int_95 = c(boot_median - crit_val_88[2]*std_err, boot_median + crit_val_88[2]*std_err)

print(paste("Confidence interval for 95% = (",
            round(confidence_int_95[1], 2),
            ",",
            round(confidence_int_95[2], 2),
            ")"))
```

```
## [1] "Confidence interval for 95% = ( 10.47 , 10.52 )"
```